

Writing applications for Cloud Foundry using Spring and MongoDB

Thomas Risberg, VMware

<trisberg@vmware.com> - @trisberg

Jared Rosoff, 10gen

<jsr@10gen.com> - @forjared

Outline

What we will cover:

- Introduction to Cloud Foundry
- Introduction to MongoDB
- Spring Data support for MongoDB
 - Using the MongoTemplate
 - Using the Mongo Repository support
- Why run MongoDB in the cloud?

What is Cloud Foundry?

- Cloud Foundry is a PaaS
 - The application platform will be delivered as a service in the cloud era
 - The industry calls this platform as a service (PaaS)
- PaaS makes it much easier to deploy, run and scale applications
- But PaaS solutions in the market have fatal flaws today
 - Limited in framework, application services and/or cloud support
- Cloud Foundry aim to fix that...

Characteristics of PaaS

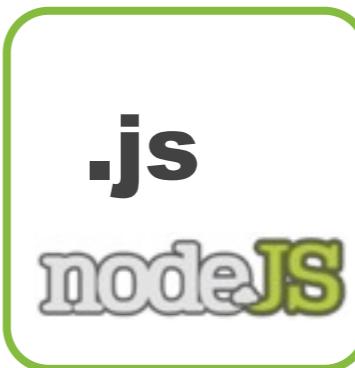
The application platform for the cloud era

- Integrated software stack
- Application execution engine
- Self-service application deployment
- Automated application infrastructure provisioning
- Curated, updated and operated as a service

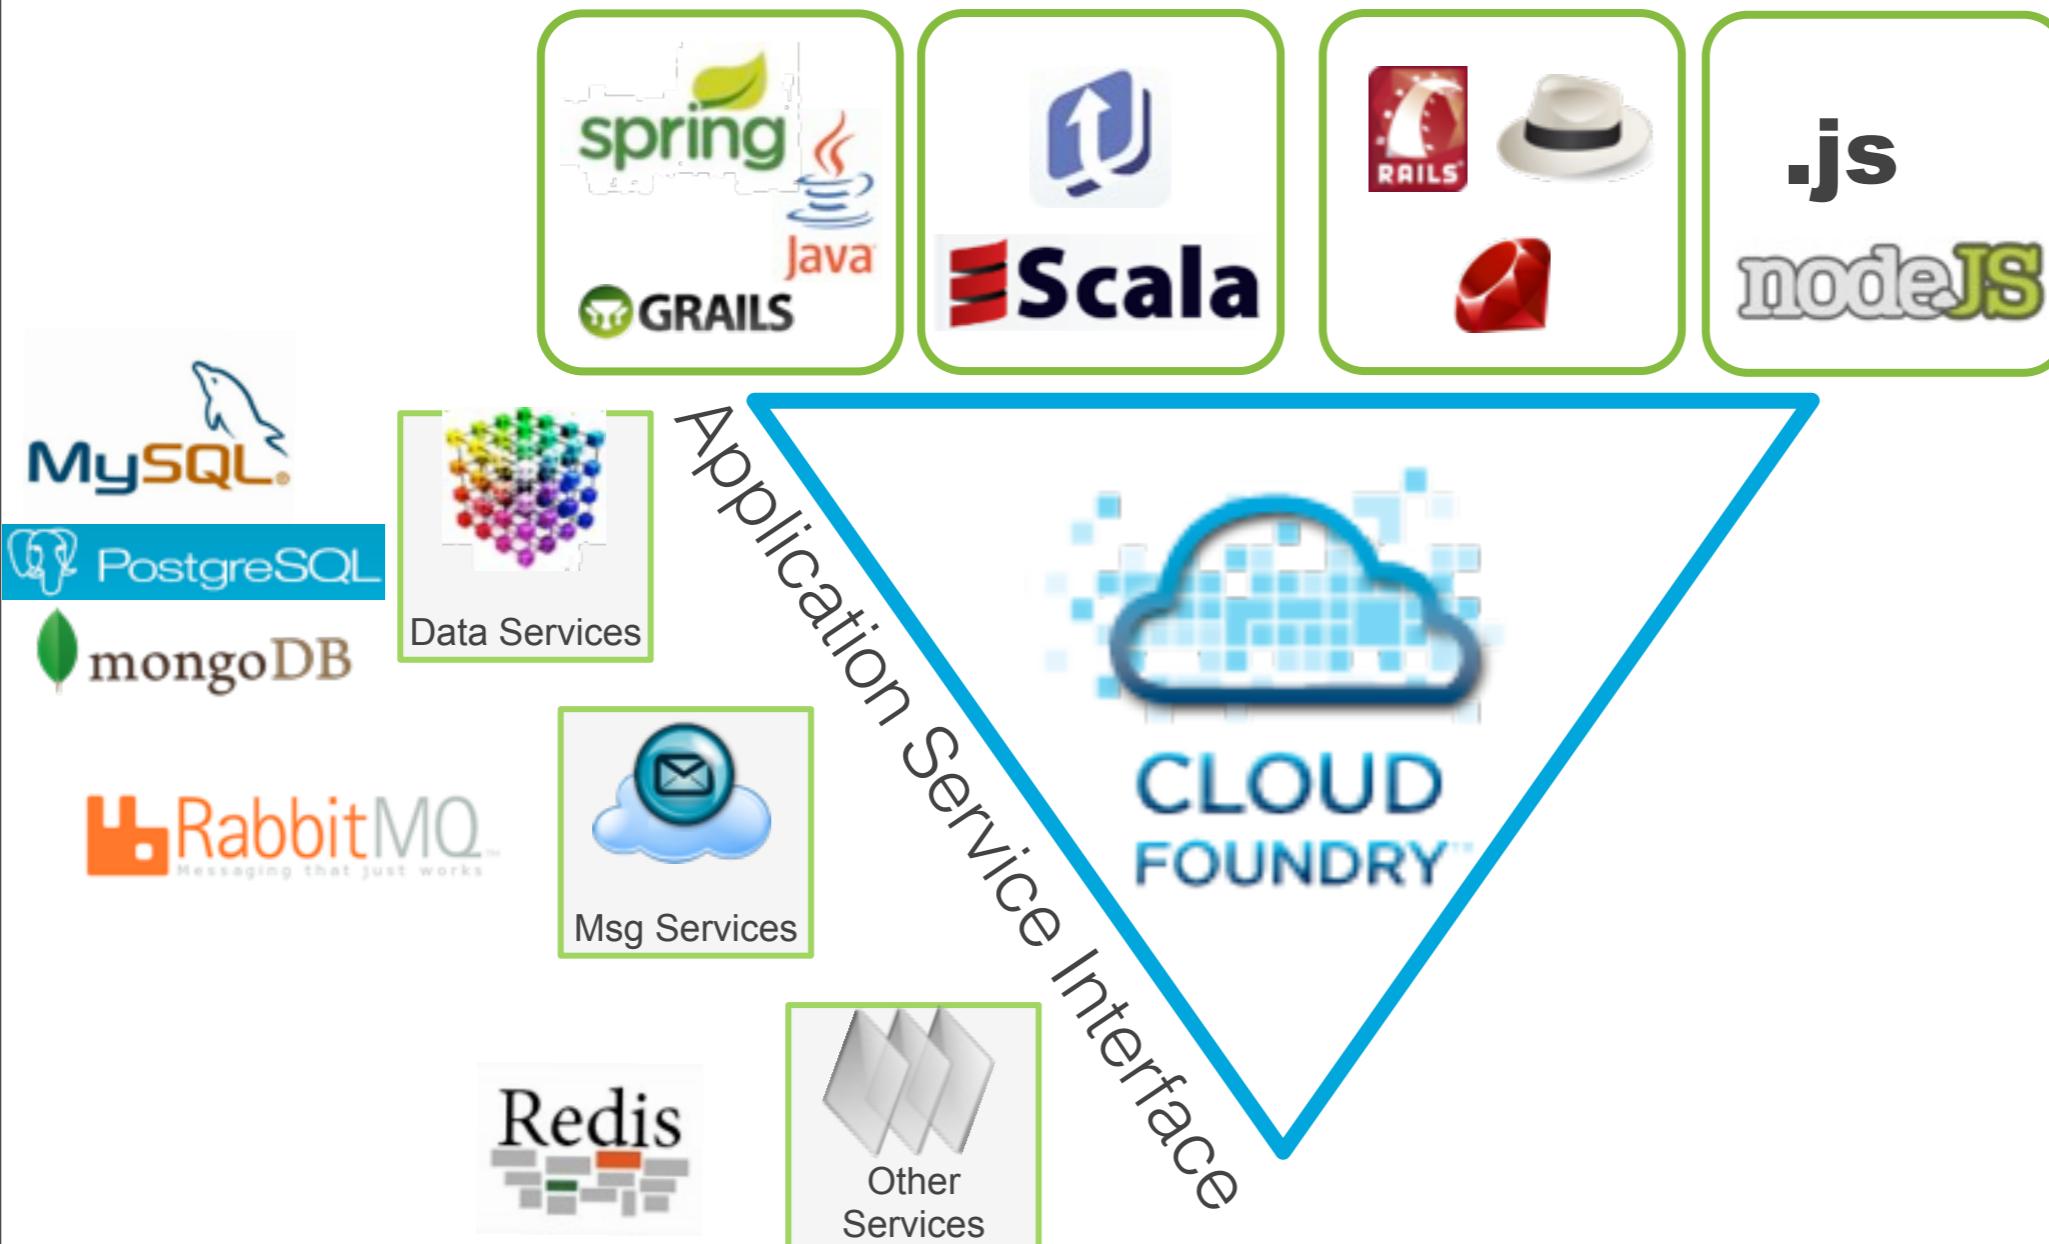
Cloud Foundry – The first open PaaS

- **Self-service application execution engine**
 - Build applications with latest high productivity frameworks
- **Automation engine for deployment and lifecycle management**
 - Deploy and cloud-scale applications in seconds
- **Open architecture**
 - Choice of clouds for deployment
 - Choice of industry-standard frameworks
 - Choice of application infrastructure services
 - Extensible architecture to “digest” future cloud innovation
 - Available as open source

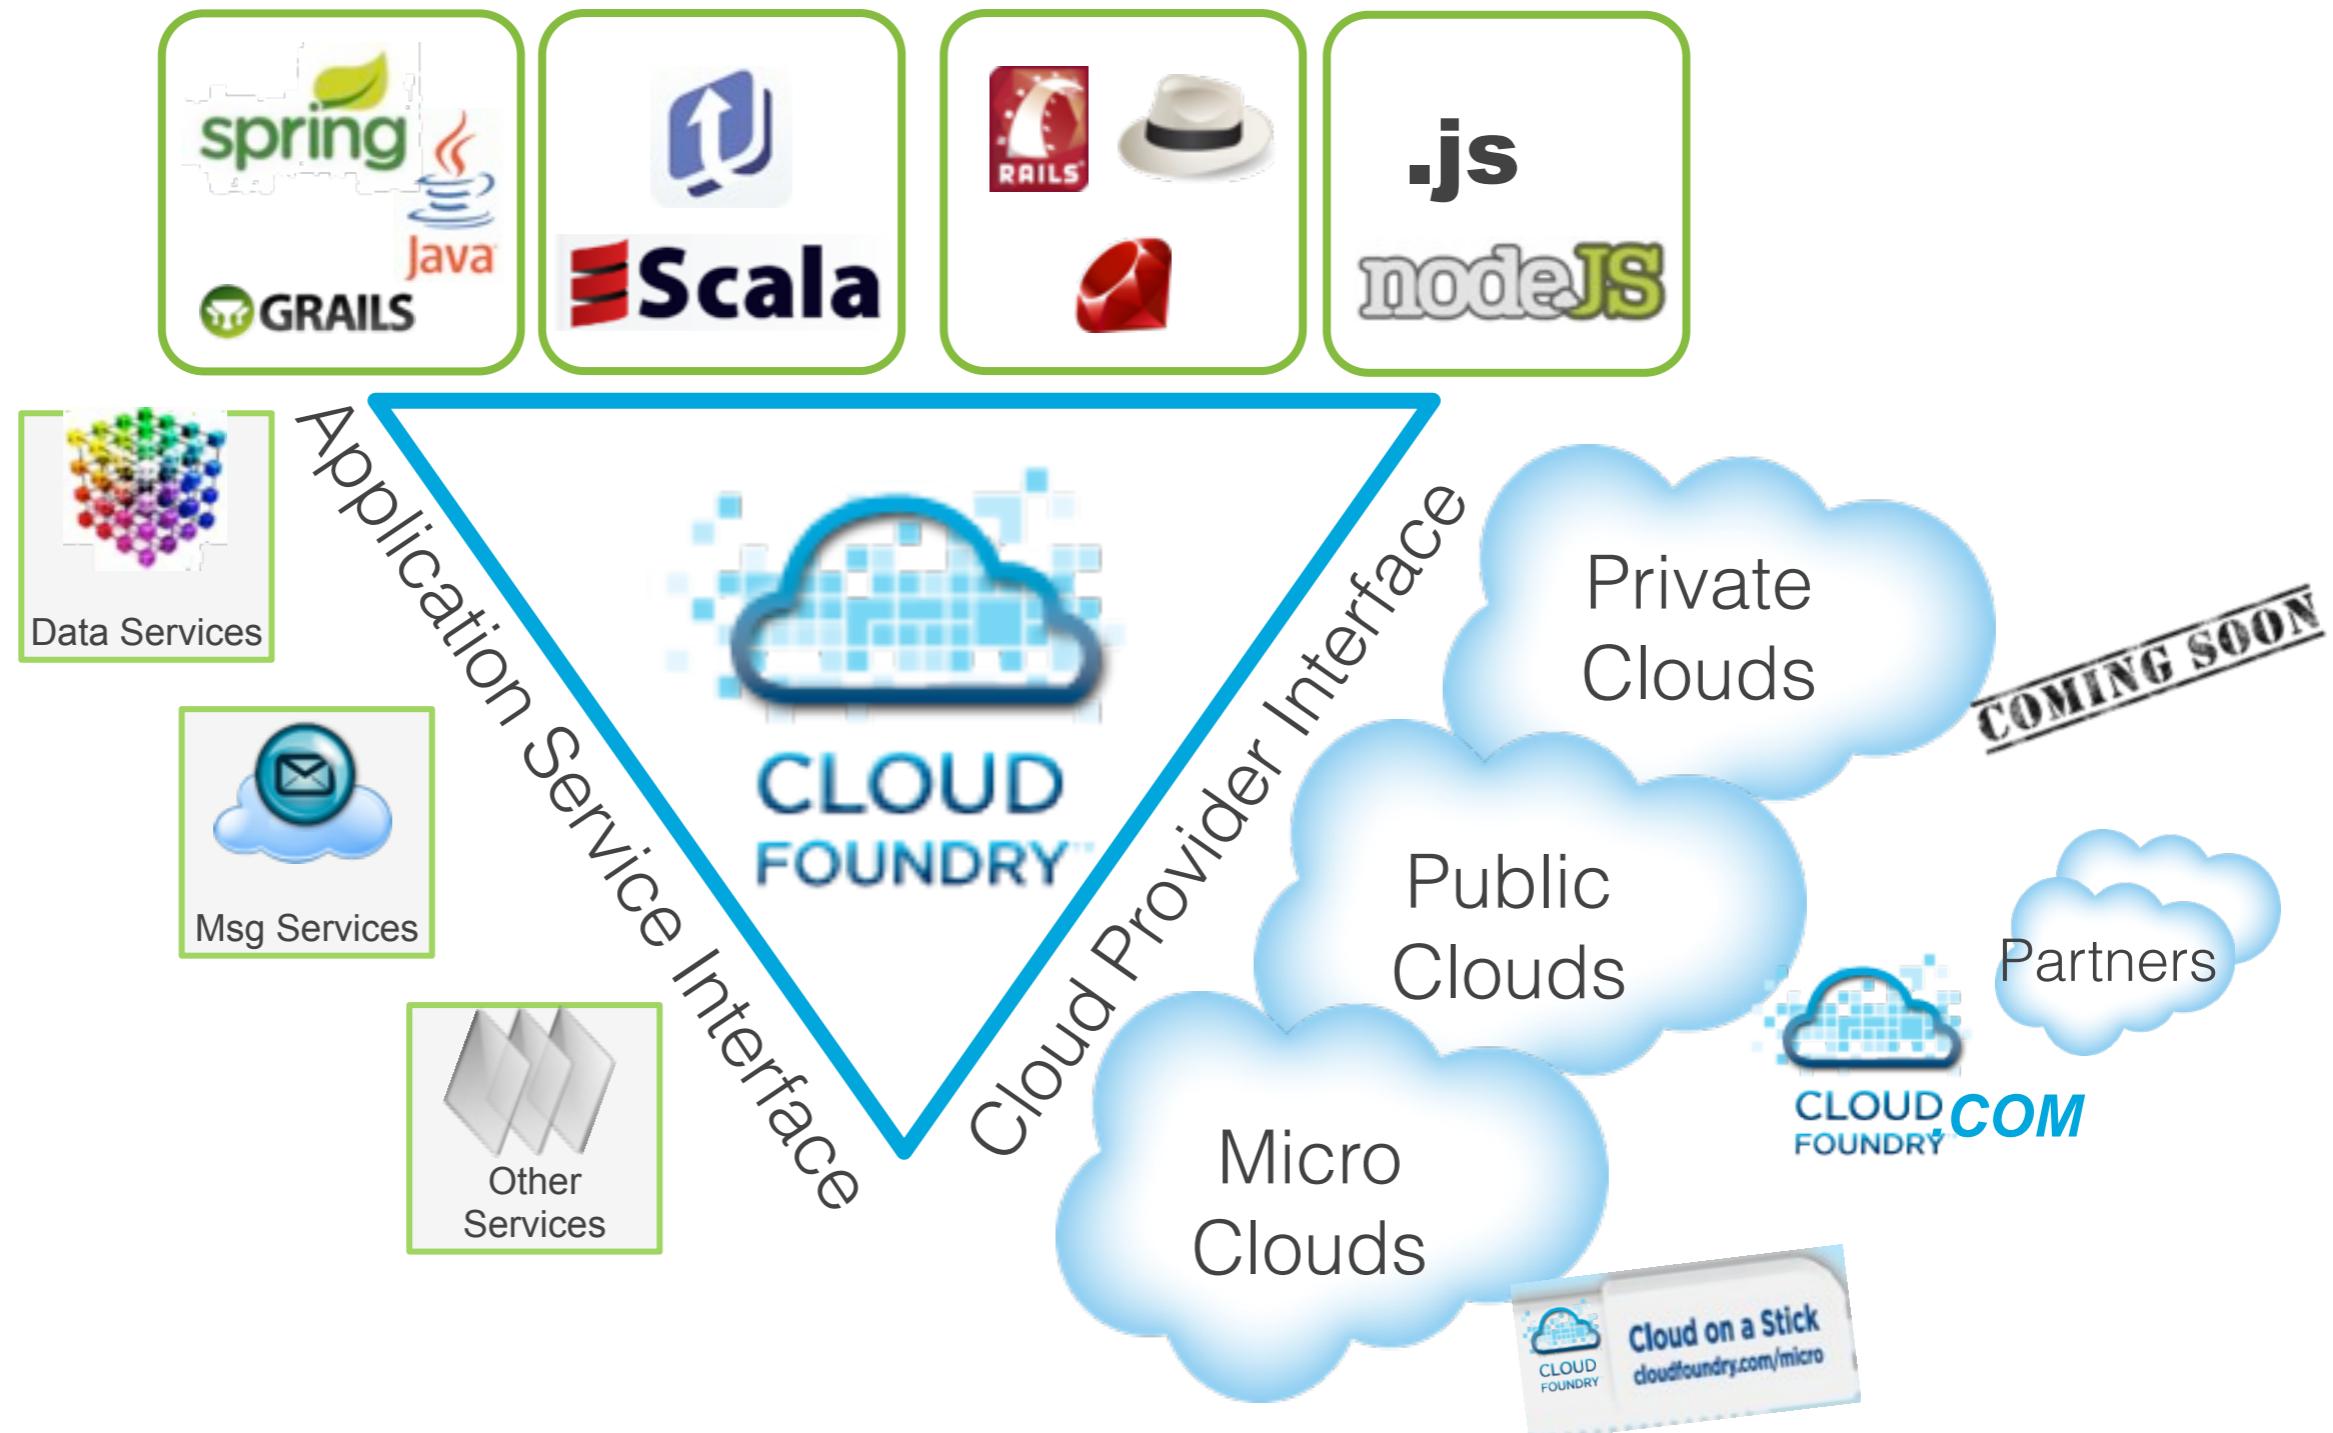
Choice of frameworks

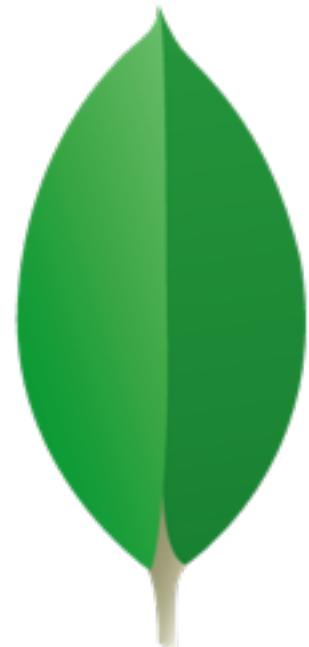


Choice of application services



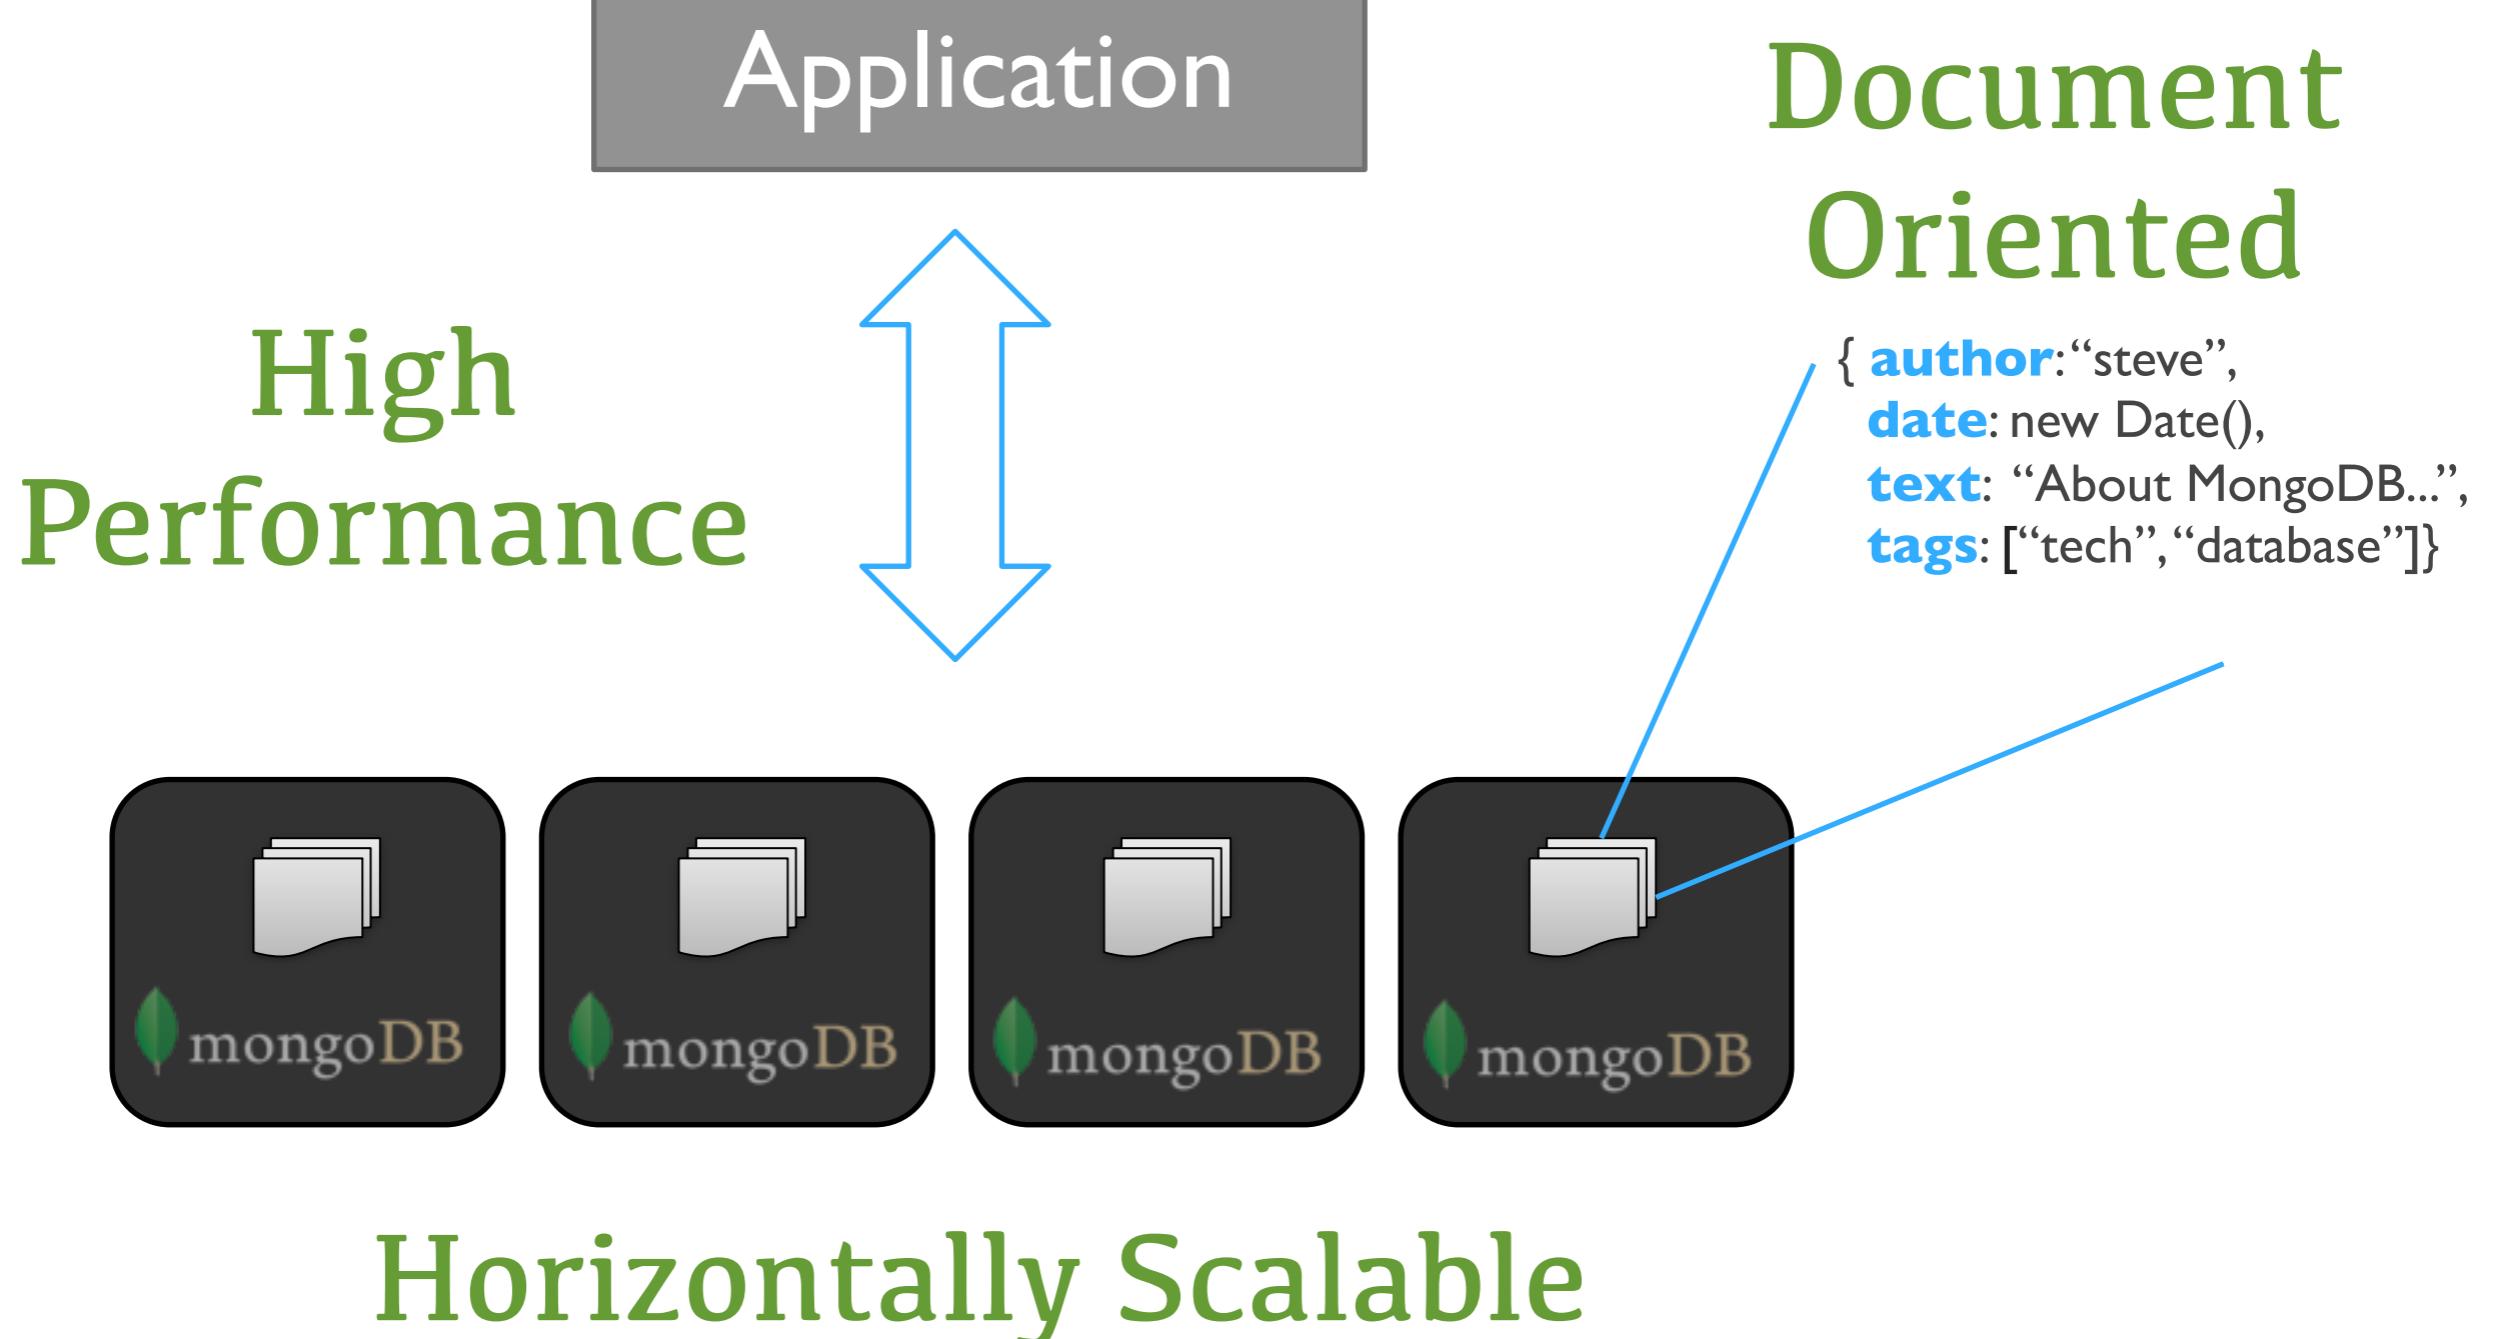
Choice of clouds





mongoDB

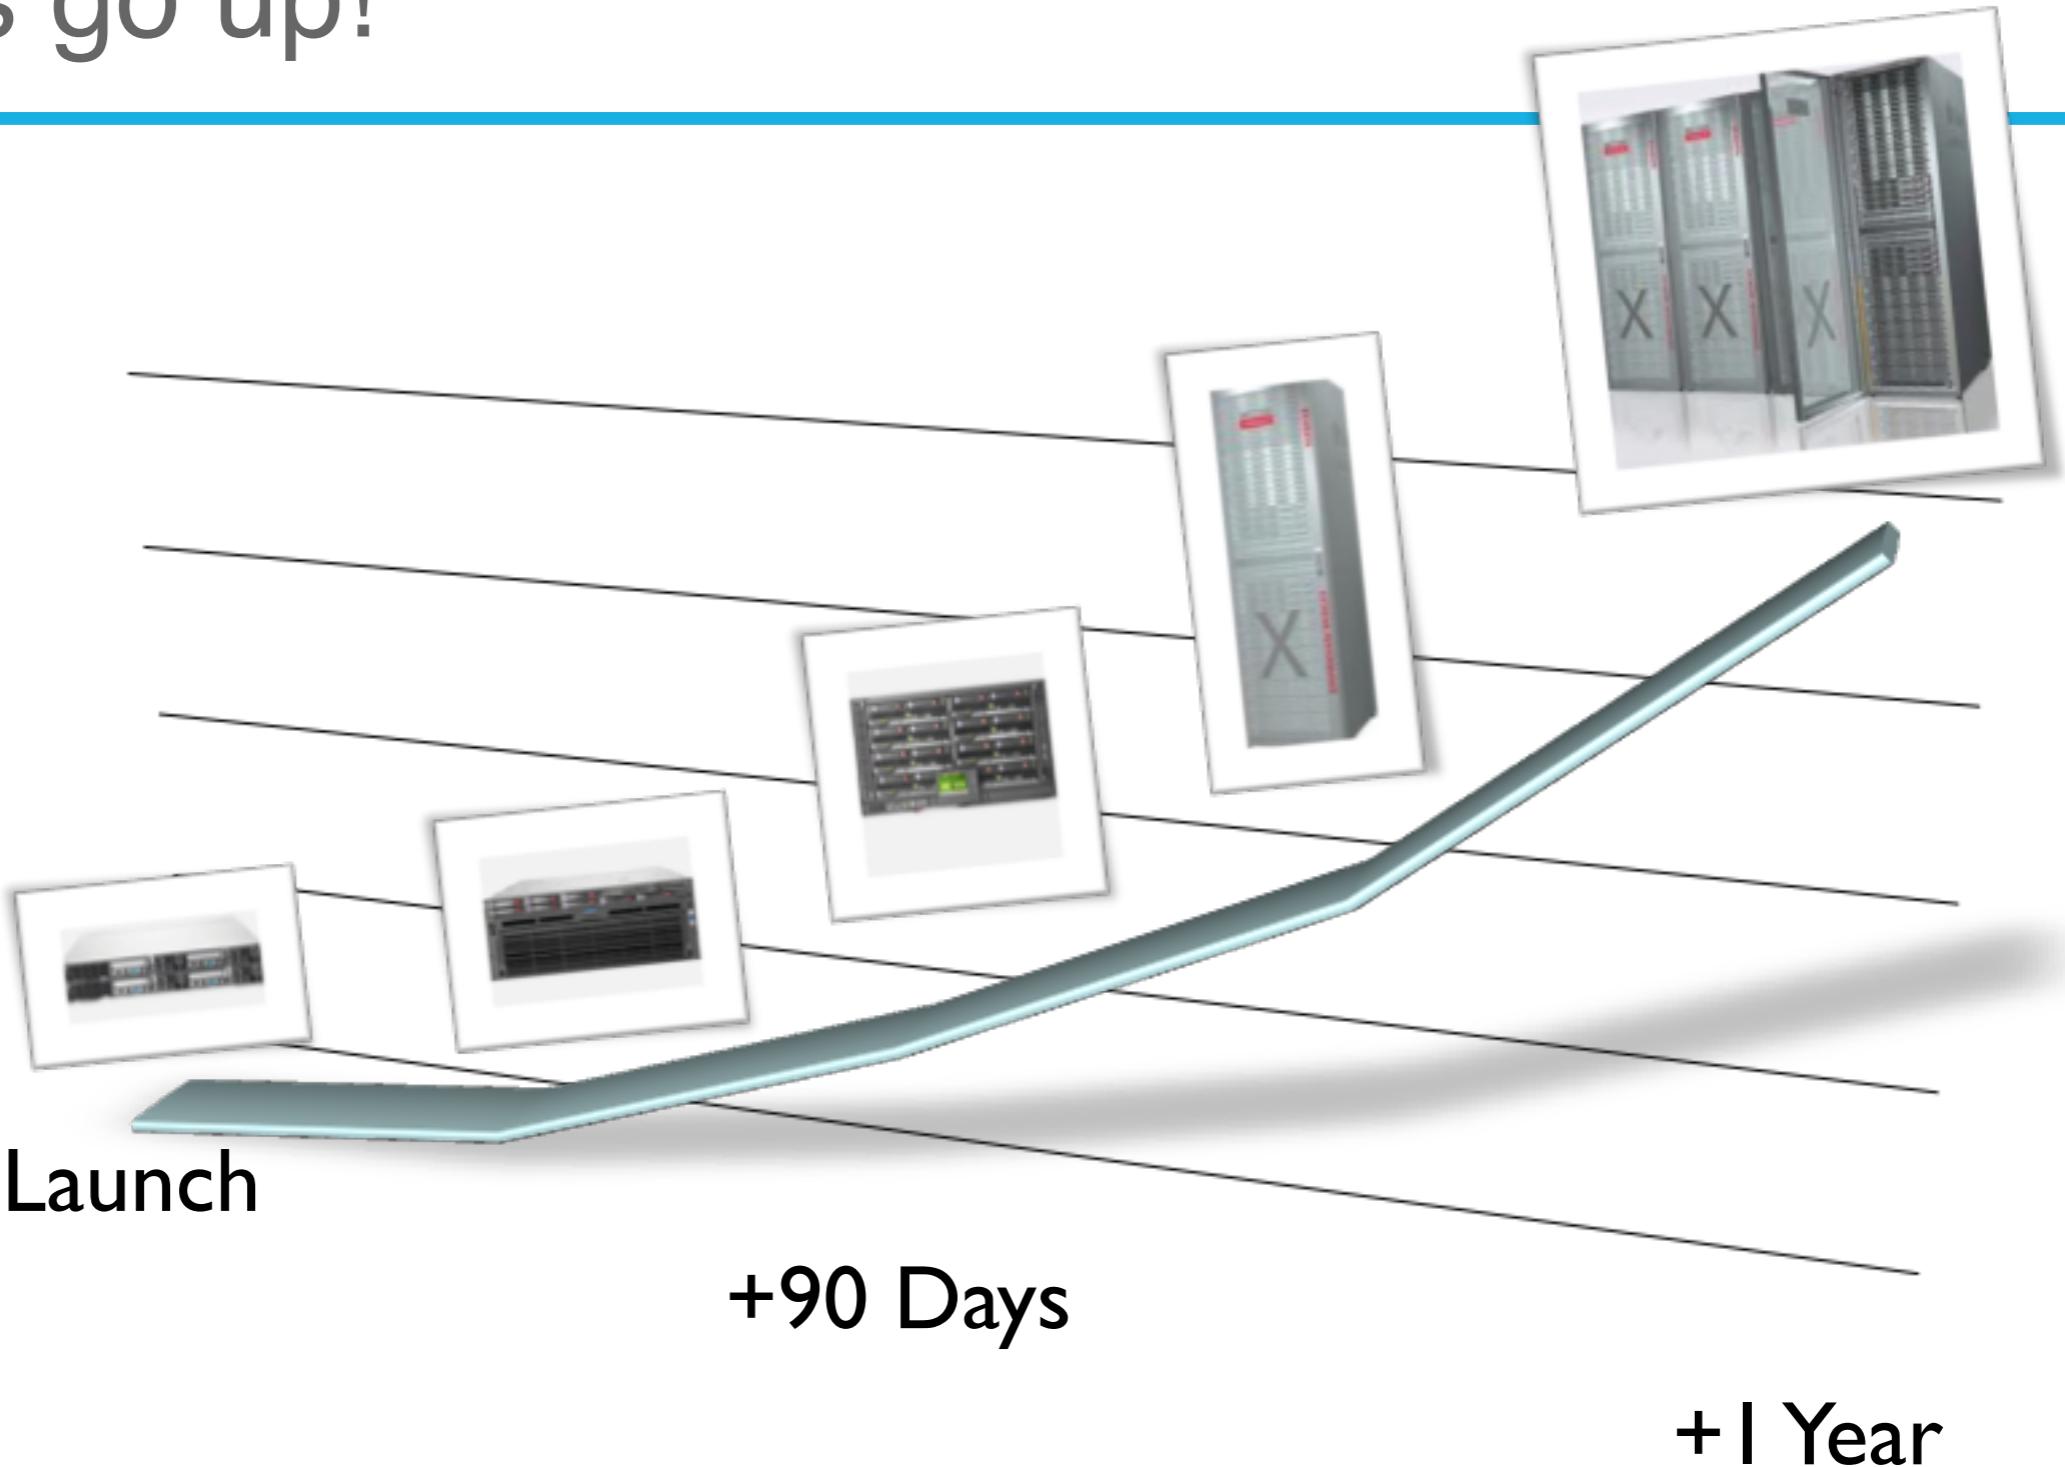
MongoDB



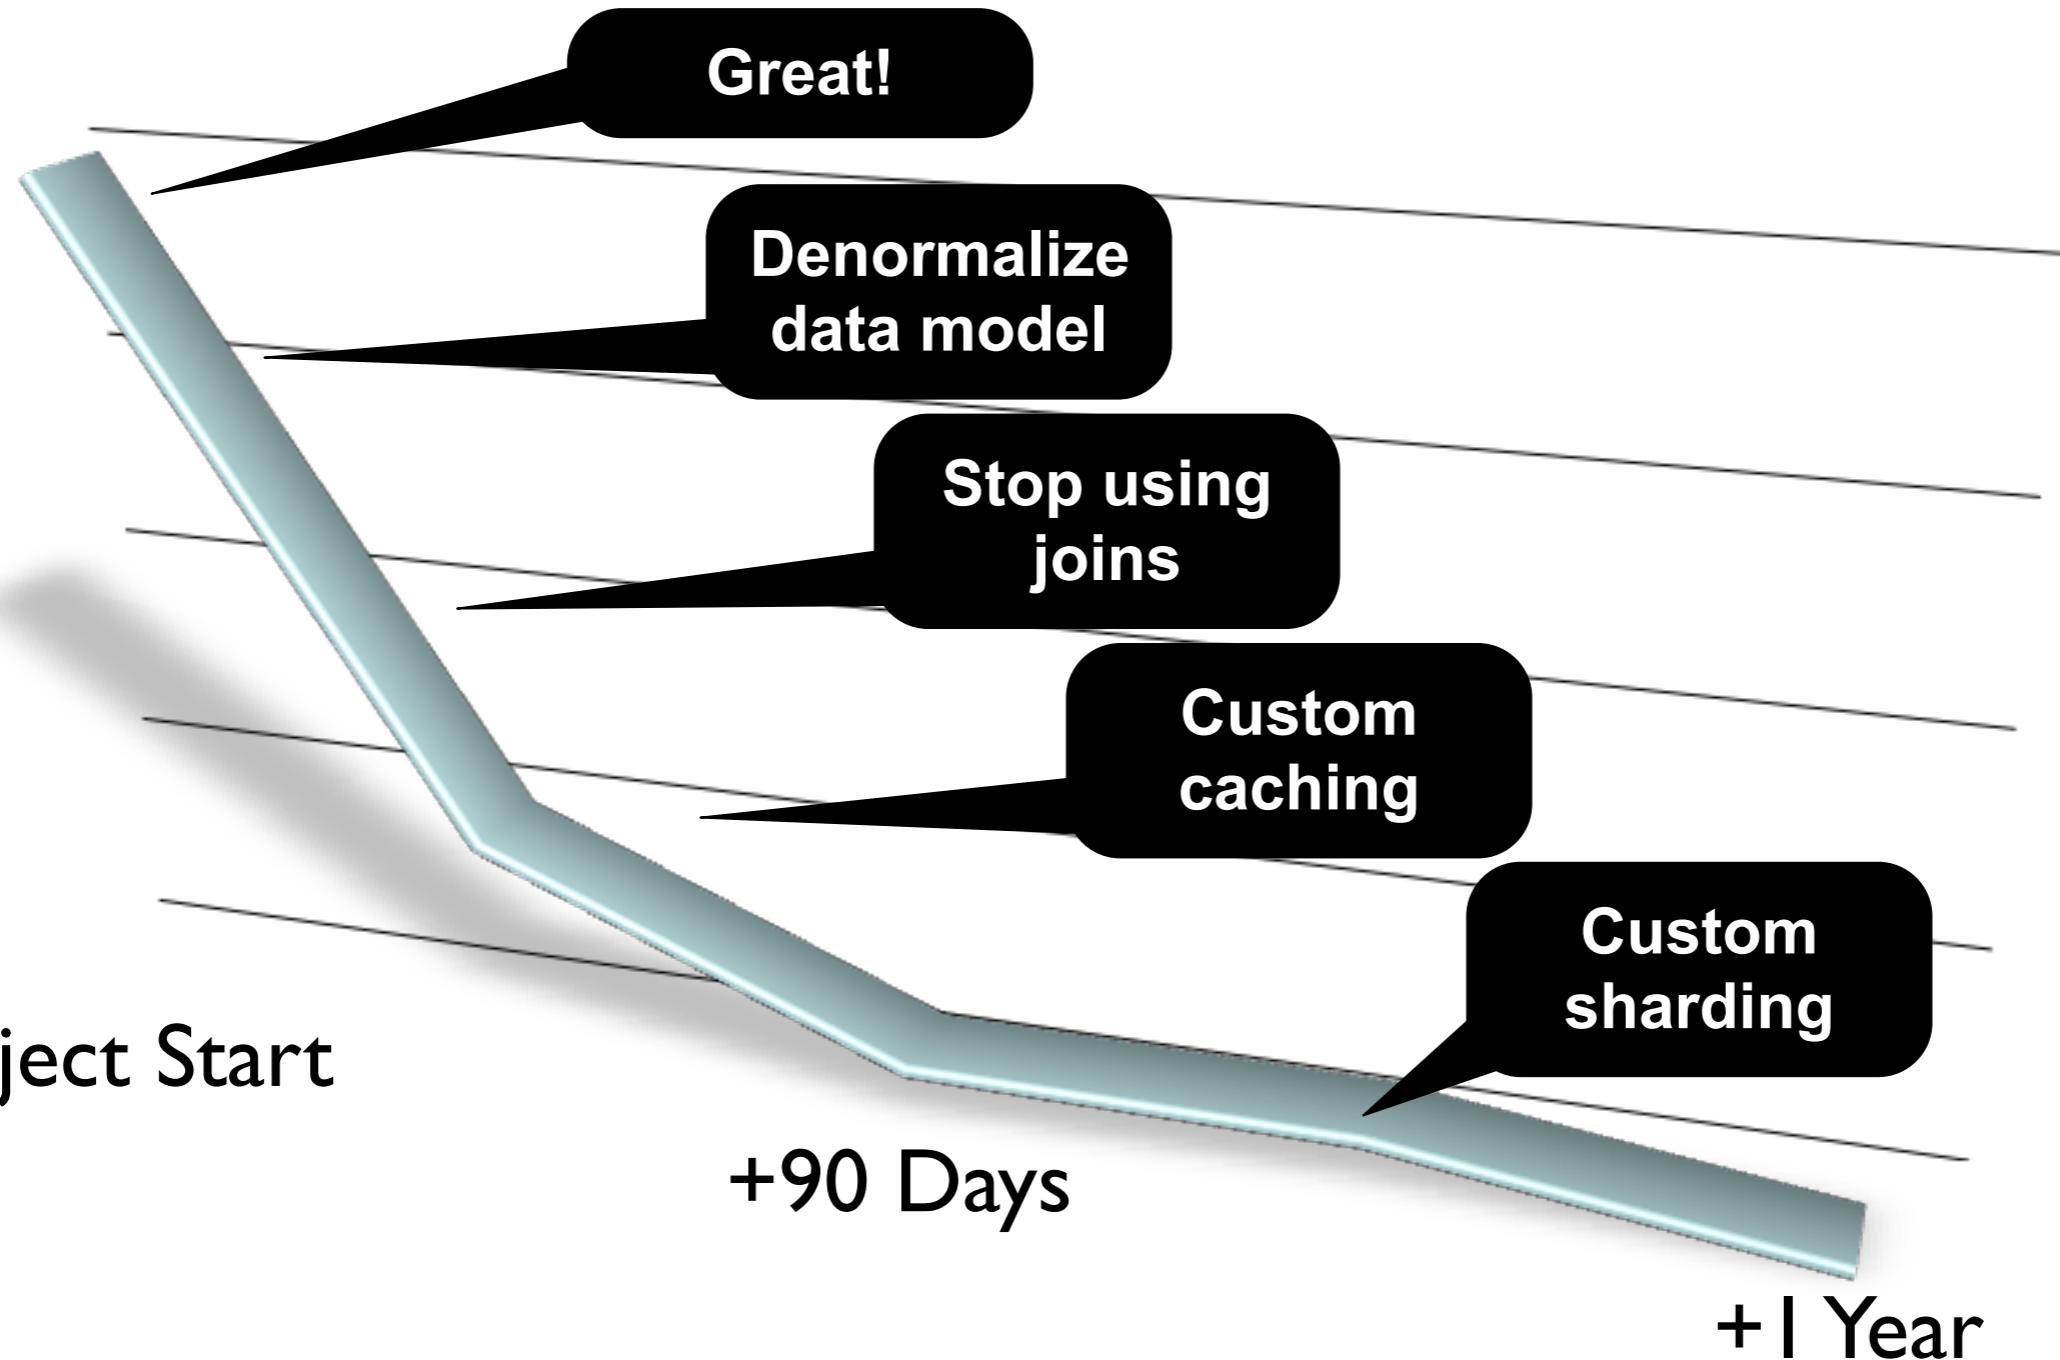
Why MongoDB

- Double click serving 400,000 ads / second
- People writing their own datastores
- Caching is de rigueur
- Complex ORM frameworks
- Computer architecture trends
- Cloud computing

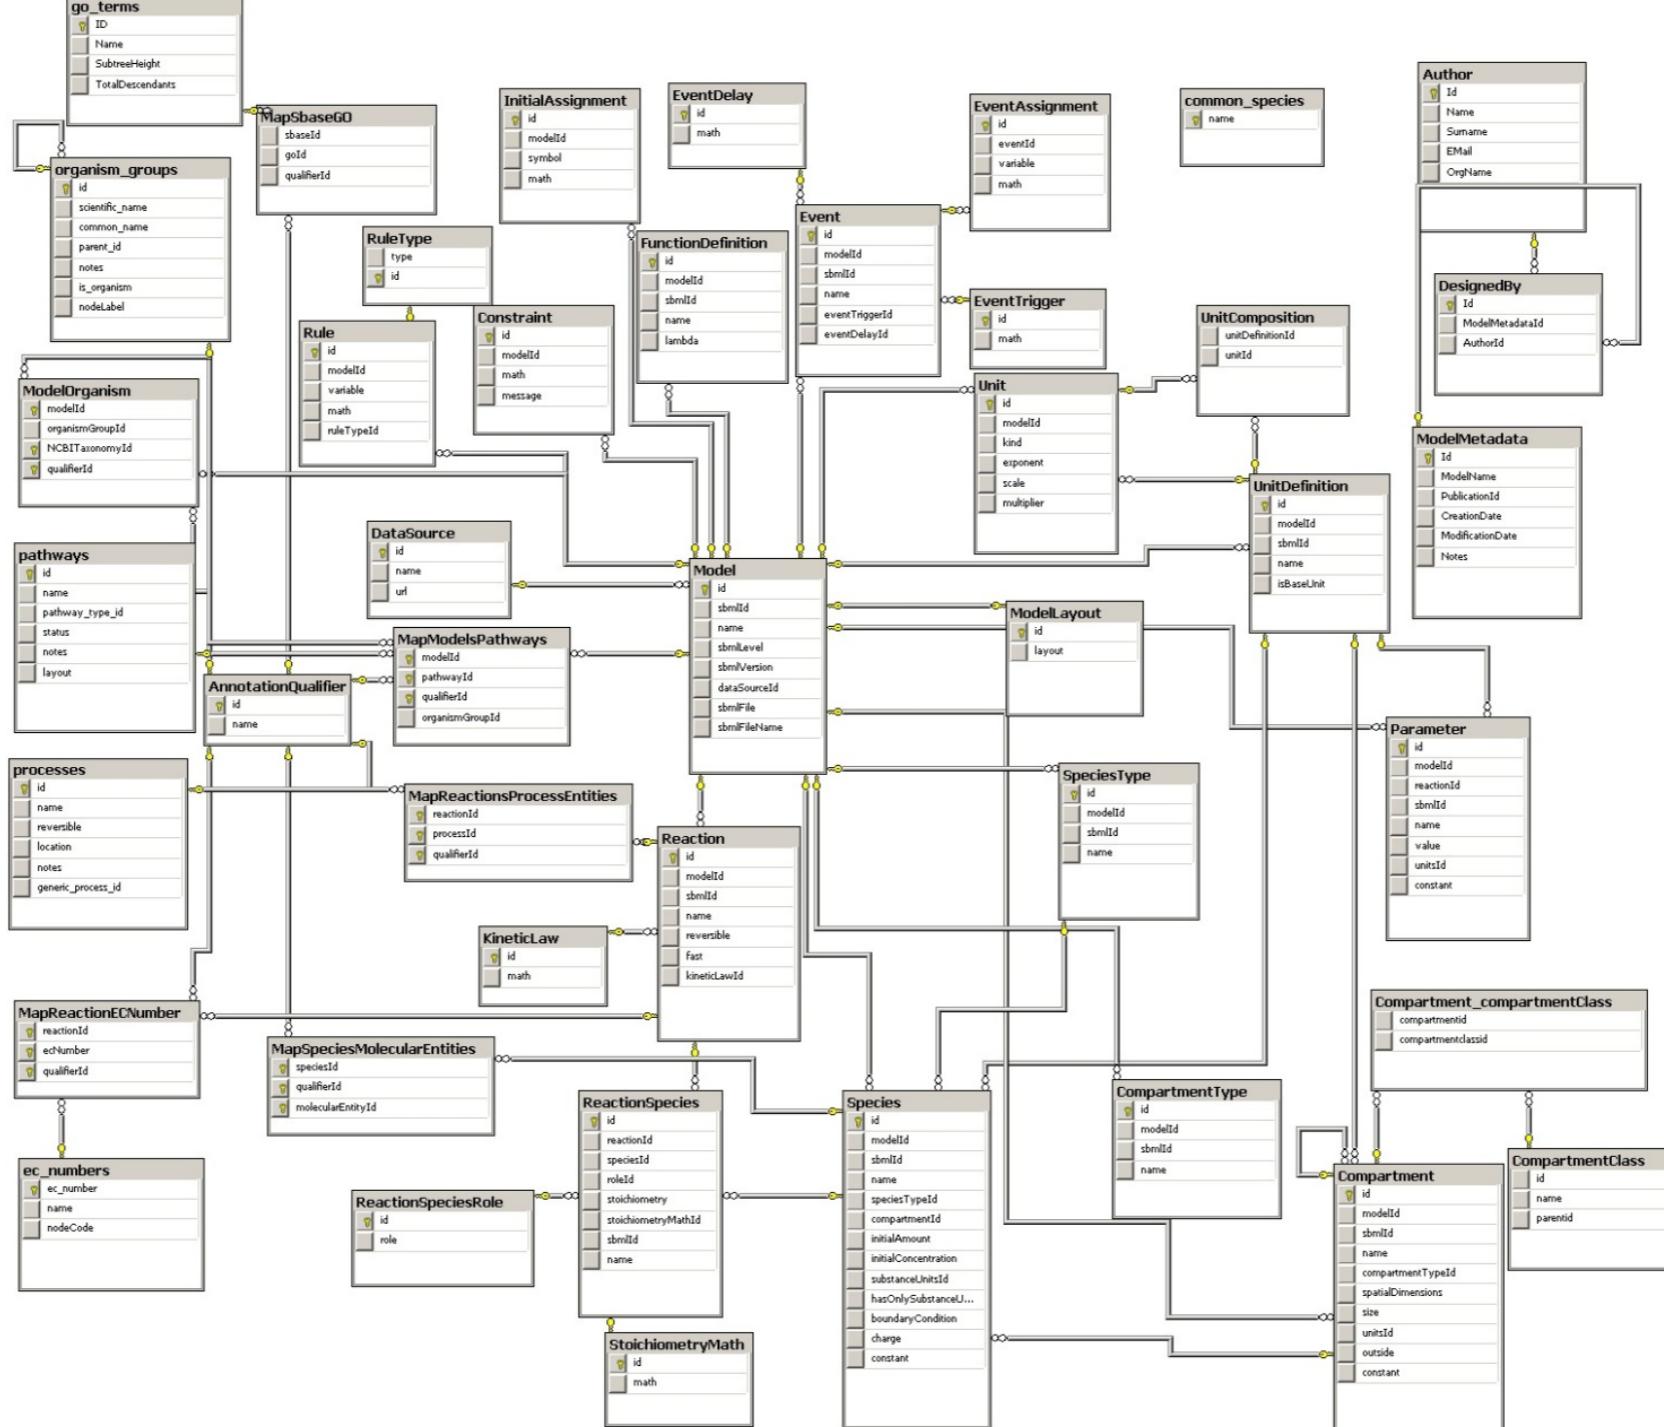
Costs go up!



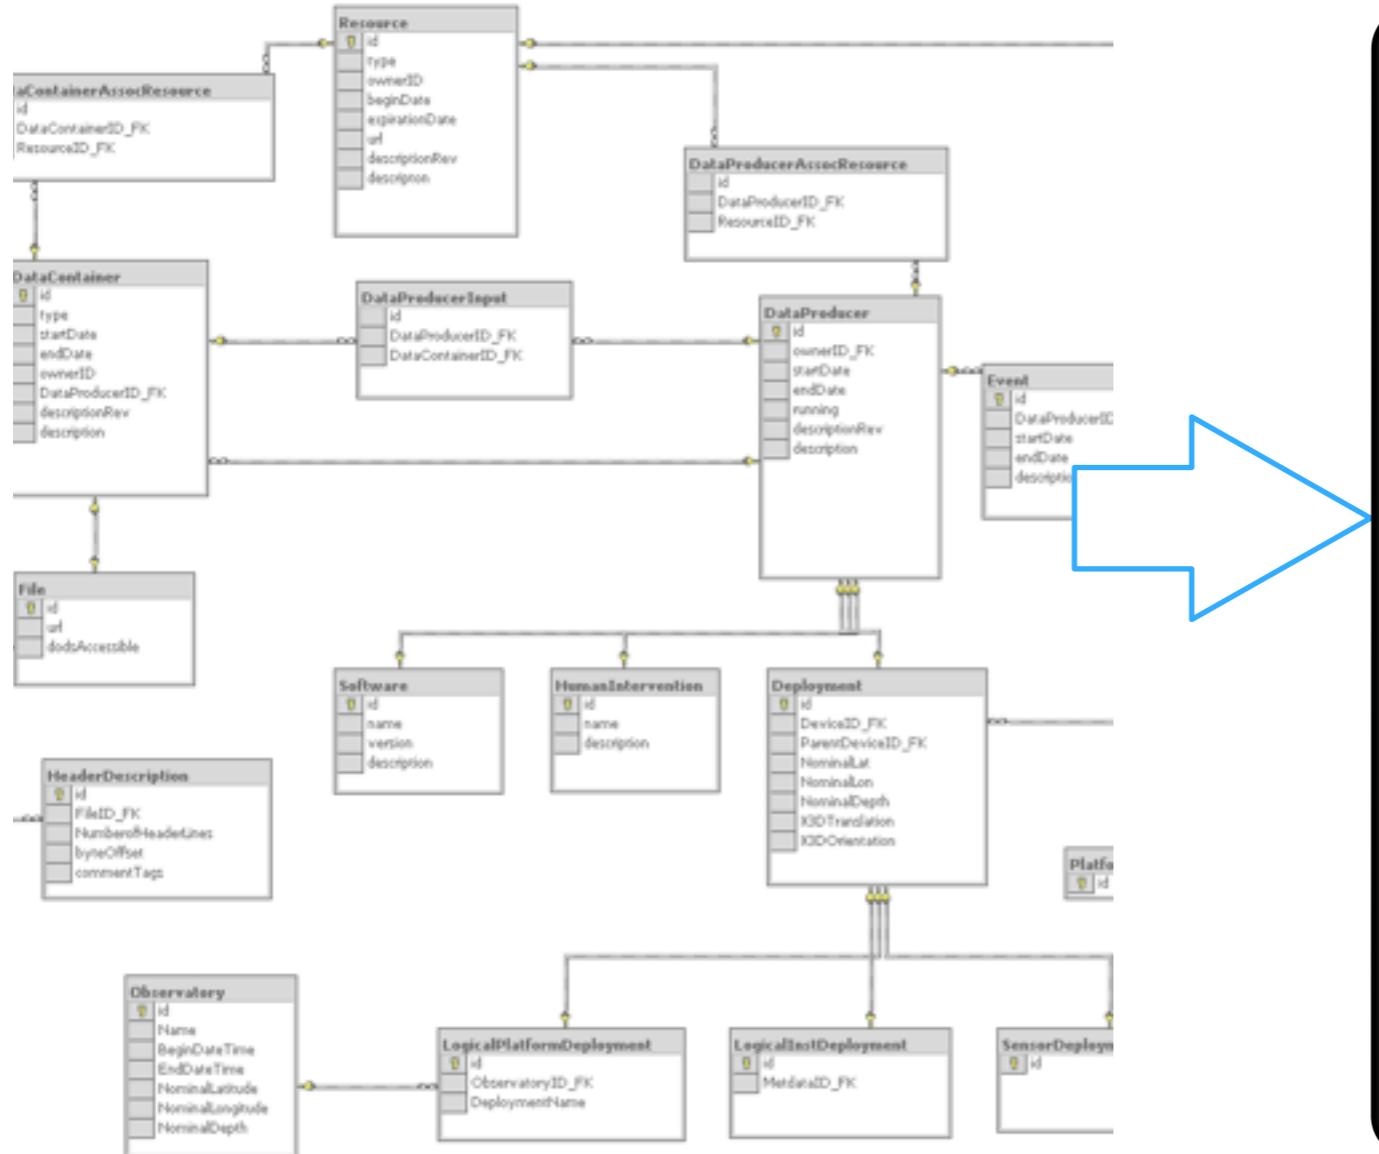
Productivity



Data models



From tables to Documents



{

```
title: 'MongoDB',
contributors: [
  { name: 'Eliot Horowitz',
    email: 'eh@10gen.com' },
  { name: 'Dwight Merriman',
    email: 'dm@10gen.com' }
```

],

model: {

```
  relational: false,
  awesome: true
```

}

}

OME zone: fine



Documents

```
> p = {author: "roger",
       date: new Date(),
       text: "about mongoDB...",
       tags: ["tech", "databases"]}

> db.posts.save(p)
```

Querying

```
> db.posts.find()  
  
> { _id : ObjectId("4c4ba5c0672c685e5e8abf3"),  
     author : "roger",  
     date : "Sat Jul 24 2010 19:47:11",  
     text : "About MongoDB...",  
     tags : [ "tech", "databases" ] }
```

Secondary Indexes

```
// 1 means ascending, -1 means descending
> db.posts.ensureIndex({author: 1})
> db.posts.find({author: 'roger'})

> { _id : ObjectId("4c4ba5c0672c685e5e8aabf3"),
  author : "roger",
  ... }
```

Conditional Query Operators

**\$all, \$exists, \$mod, \$ne, \$in, \$nin, \$nor,
\$or, \$size, \$type, \$lt, \$lte, \$gt, \$gte**

```
// find posts with any tags  
> db.posts.find( {tags: {$exists: true}} )
```

```
// find posts matching a regular expression  
> db.posts.find( {author: /^rog*/i } )
```

```
// count posts by author  
> db.posts.find( {author: 'roger'} ).count()
```

```
> qp.posts.find( {author: 'roger'} ).count()  
\\\ count posts by author
```

Atomic Updates

**\$set, \$unset, \$inc, \$push, \$pushAll,
\$pull, \$pullAll, \$bit**

```
> comment = { author: "fred",
    date: new Date(),
    text: "Best Movie Ever"}  
  
> db.posts.update( { _id: "..."},  
    $push: {comments: comment} );
```

```
    $push: {comments: comment} );  
> db.posts.update( { _id: "..."},
```

Nested Documents

```
{ _id : ObjectId("4c4ba5c0672c685e5e8aabf3"),
  author : "roger",
  date : "Sat Apr 24 2011 19:47:11",
  text : "About MongoDB...",
  tags : [ "tech", "databases" ],
  comments : [
    {
      author : "Fred",
      date : "Sat Apr 25 2010 20:51:03 GMT-0700",
      text : "Best Post Ever!"
    }
  ]
}
```

```
}
```

```
]
```

Advanced indexing

```
// Index nested documents
> db.posts.ensureIndex( "comments.author": 1)
> db.posts.find({'comments.author': 'Fred'})  
  
// Index on tags (multi-key index)
> db.posts.ensureIndex( tags: 1)
> db.posts.find( { tags: 'tech' } )  
  
// geospatial index
> db.posts.ensureIndex( "author.location": "2d" )
> db.posts.find( "author.location": { $near : [22,42] } )
```

```
> qp.posts.find( "author.location": { $near : [55,45] } )
> qp.posts.ensureIndex( "author.location": "2d" )
```

Rich Documents

```
{   _id : ObjectId("4c4ba5c0672c685e5e8abf3"),  
  
  line_items : [ { sku: 'tt-123',  
                  name: 'Coltrane: Impressions' },  
                 { sku: 'tt-457',  
                  name: 'Davis: Kind of Blue' } ],  
  
  address : { name: 'Banker',  
              street: '111 Main',  
              zip: 10010 },  
  
  payment: { cc: 4567,  
             exp: Date(2011, 7, 7) },  
  
  subtotal: 2355  
}  
  
}
```

Demo

Cloud Foundry and Mongo DB

Spring Data Intoduction



What is Spring Data all about?

- Bring classic Spring value propositions to NOSQL
 - Productivity
 - Programming model consistency
- Support for a wide range of NOSQL databases
- Also, support for new features for JDBC and JPA
- Support for other data related projects like Hadoop and Gemfire

Spring Data MongoDB Support

Spring Data support for MongoDB:

- MongoTemplate
 - ✓ MongoConverter interface for mapping Mongo documents
 - Built-in Advanced Mapping
 - Annotation based (@Document, @Id, @DbRef)
 - MappingMongoConverter for POJO mapping support
 - Leverage Spring 3.0 TypeConverters and SpEL
 - ✓ Exception translation
 - ✓ Java based Query, Criteria, and Update DSLs
- MongoRepository
 - ✓ Built on Spring Data JPA (Hades) support for JPA Repositories
 - ✓ QueryDSL integration to support type-safe queries.



Working with POJOs

```
@Document  
public class Book {  
  
    private String title;  
  
    @DBRef  
    private Author author;  
  
    @Id  
    private String isbn;  
  
    private BigDecimal price;  
  
    private Date published;  
  
    private Set<String> categories;  
  
    // getters and setters  
}
```

```
@Document  
public class Author {  
  
    @Id @SuppressWarnings("unused")  
    private String id;  
  
    private String name;  
  
    public String getName() {  
        return name;  
    }  
  
    public void setName(String name) {  
        this.name = name;  
    }  
}
```

Mongo Template CRUD Methods

```
Mongo mongo = new Mongo("localhost", 27017);
MongoDbFactory mongoDbFactory =
    new SimpleMongoDbFactory(mongo, "db");
MongoTemplate mongoTemplate = new
MongoTemplate(mongoDbFactory);

mongoTemplate.insert(book);

List<Book> books = mongoTemplate.findAll(Book.class);

Book book = mongoTemplate.findOne(query, Book.class);

mongoTemplate.save(book);

mongoTemplate.remove(query, Book.class);
```

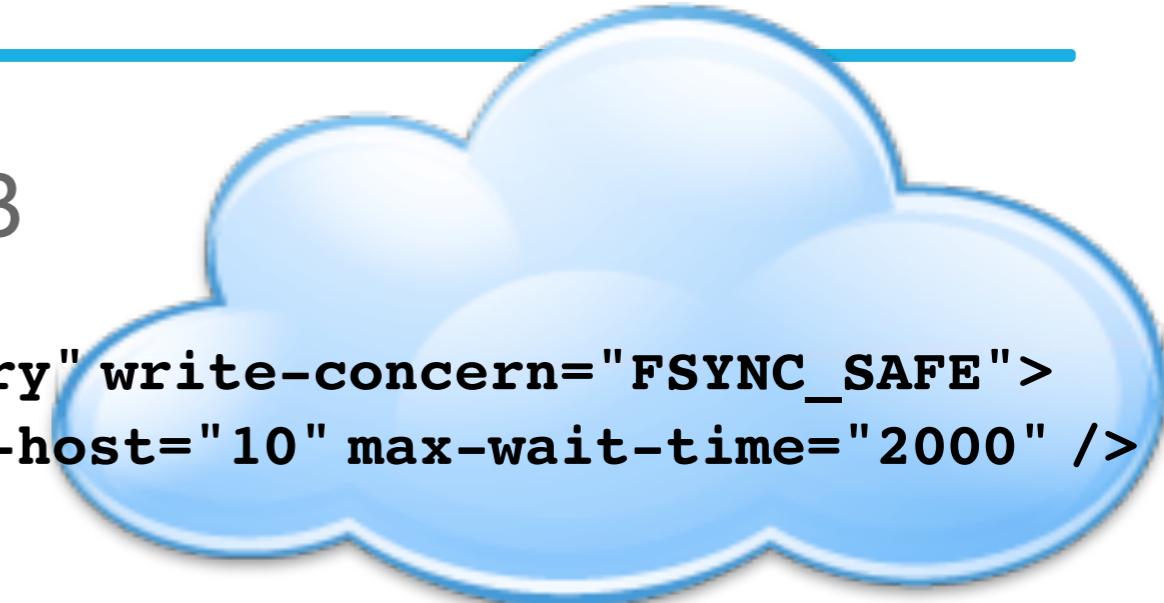
Querying

```
public List<Book> findByCategoriesOrYear(Set<String> categories, String year) {
    String[] categoriesToMatch;
    if (categories == null) {
        categoriesToMatch = new String[] {};
    }
    else {
        categoriesToMatch = categories.toArray(new String[categories.size()]);
    }
    Date startDate = null;
    if (year != null && year.length() == 4) {
        DateFormat formatter = new SimpleDateFormat("yyyy-dd-MM");
        try {
            startDate = formatter.parse(year + "-01-01");
        } catch (ParseException e) {}
    }
    Criteria searchCriteria = null;
    if (startDate != null) {
        searchCriteria =
            Criteria.where("published").gte(startDate);
    }
    else {
        searchCriteria = new Criteria();
    }
    if (categoriesToMatch.length > 0) {
        searchCriteria
            .and("categories").in((Object[])categoriesToMatch);
    }
    Query query = new Query(searchCriteria);
    return mongoTemplate.find(query, Book.class);
}
```

Cloud Foundry MongoDB Support

Namespace support for MongoDB

```
<cloud:mongo-db-factory id="mongoDbFactory" write-concern="FSYNC_SAFE">
  <cloud:mongo-options connections-per-host="10" max-wait-time="2000" />
</cloud:mongo-db-factory>
```



```
<dependency>
  <groupId>org.cloudfoundry</groupId>
  <artifactId>cloudfoundry-runtime</artifactId>
  <version>0.8.1</version>
</dependency>

<repository>
  <id>org.springframework.maven.milestone</id>
  <name>Spring Framework Maven Milestone Repository</name>
  <url>http://maven.springframework.org/milestone</url>
</repository>
```

Demo

MongoTemplate example:

<https://github.com/trisberg/springone-mongotemplate>

Spring Data Repository

Spring Data Repository basics:

- Generic repository implementation
- Basic CRUD (create, read, update and delete) methods
- Generating code for queries defined in repository interface
 - findAll
 - findByName ...
- Pagination and sorting support
- Currently has JPA and Mongo implementations

Spring Data Repository Interfaces

CrudRepository

long	<code>count()</code> Returns the number of entities available.
void	<code>delete(ID id)</code> Deletes the entity with the given id.
void	<code>delete(Iterable<? extends T> entities)</code> Deletes the given entities.
void	<code>delete(T entity)</code> Deletes a given entity.
void	<code>deleteAll()</code> Deletes all entities managed by the repository.
boolean	<code>exists(ID id)</code> Returns whether an entity with the given id exists.
Iterable<T>	<code>findAll()</code> Returns all instances of the type.
T	<code>findOne(ID id)</code> Retrieves an entity by its primary key.
Iterable<T>	<code>save(Iterable<? extends T> entities)</code> Saves all given entities.
T	<code>save(T entity)</code> Saves a given entity.

PagingAndSortingRepository

<code>Page<T></code>	<code>findAll(Pageable pageable)</code> Returns a <code>Page</code> of entities meeting the paging restriction provided in the <code>Pageable</code> object.
<code>Iterable<T></code>	<code>findAll(Sort sort)</code> Returns all entities sorted by the given options.

Spring Data Repository Query Methods

Spring Data Repository query methods:

Keyword	Sample	Logical result
GreaterThan	findByAgeGreaterThan(int age)	{"age" : {"\$gt" : age}}
LessThan	findByAgeLessThan(int age)	{"age" : {"\$lt" : age}}
Between	findByAgeBetween(int from, int to)	{"age" : {"\$gt" : from, "\$lt" : to}}
IsNotNull, NotNull	findByNameNotNull()	{"name" : {"\$ne" : null}}
IsNull, Null	findByNameNull()	{"name" : null}
Like	findByNameLike(String expr)	{"name" : expr} (expr as regex)
(No keyword)	findByName(String name)	{"name" : name}
Not	findByNameNot(String name)	{"name" : {"\$ne" : name}}

and more, see reference docs ...

Mongo Repository Example

```
public interface BookRepository extends Repository<Book, String> {  
    Book save(Book book);  
  
    Book findOne(String isbn);  
  
    void delete(String isbn);  
  
    List<Book> findAll();  
  
    List<Book> findByPublishedGreaterThanOrEqual(Date date);  
  
    List<Book> findByCategoriesIn(String[] categories);  
  
    List<Book> findByPublishedGreaterThanOrEqualAndCategoriesIn(  
        Date date, String[] categories);  
}
```

Demo

Mongo Repository example:

<https://github.com/trisberg/springone-repository>

QueryDSL Repository

QueryDSL allows you to build typesafe queries using Java for JPA, SQL, MongoDB and more

- Generic repository implementation
- Code completion in IDE
- No Strings
 - reference types and properties safely
 - easier refactoring
- Incremental query definition is easier
- Annotation preprocessor generates query types

Demo

Mongo QueryDSL Repository example:

<https://github.com/trisberg/springone-repo-qdsl>

Why run MongoDB in the cloud?



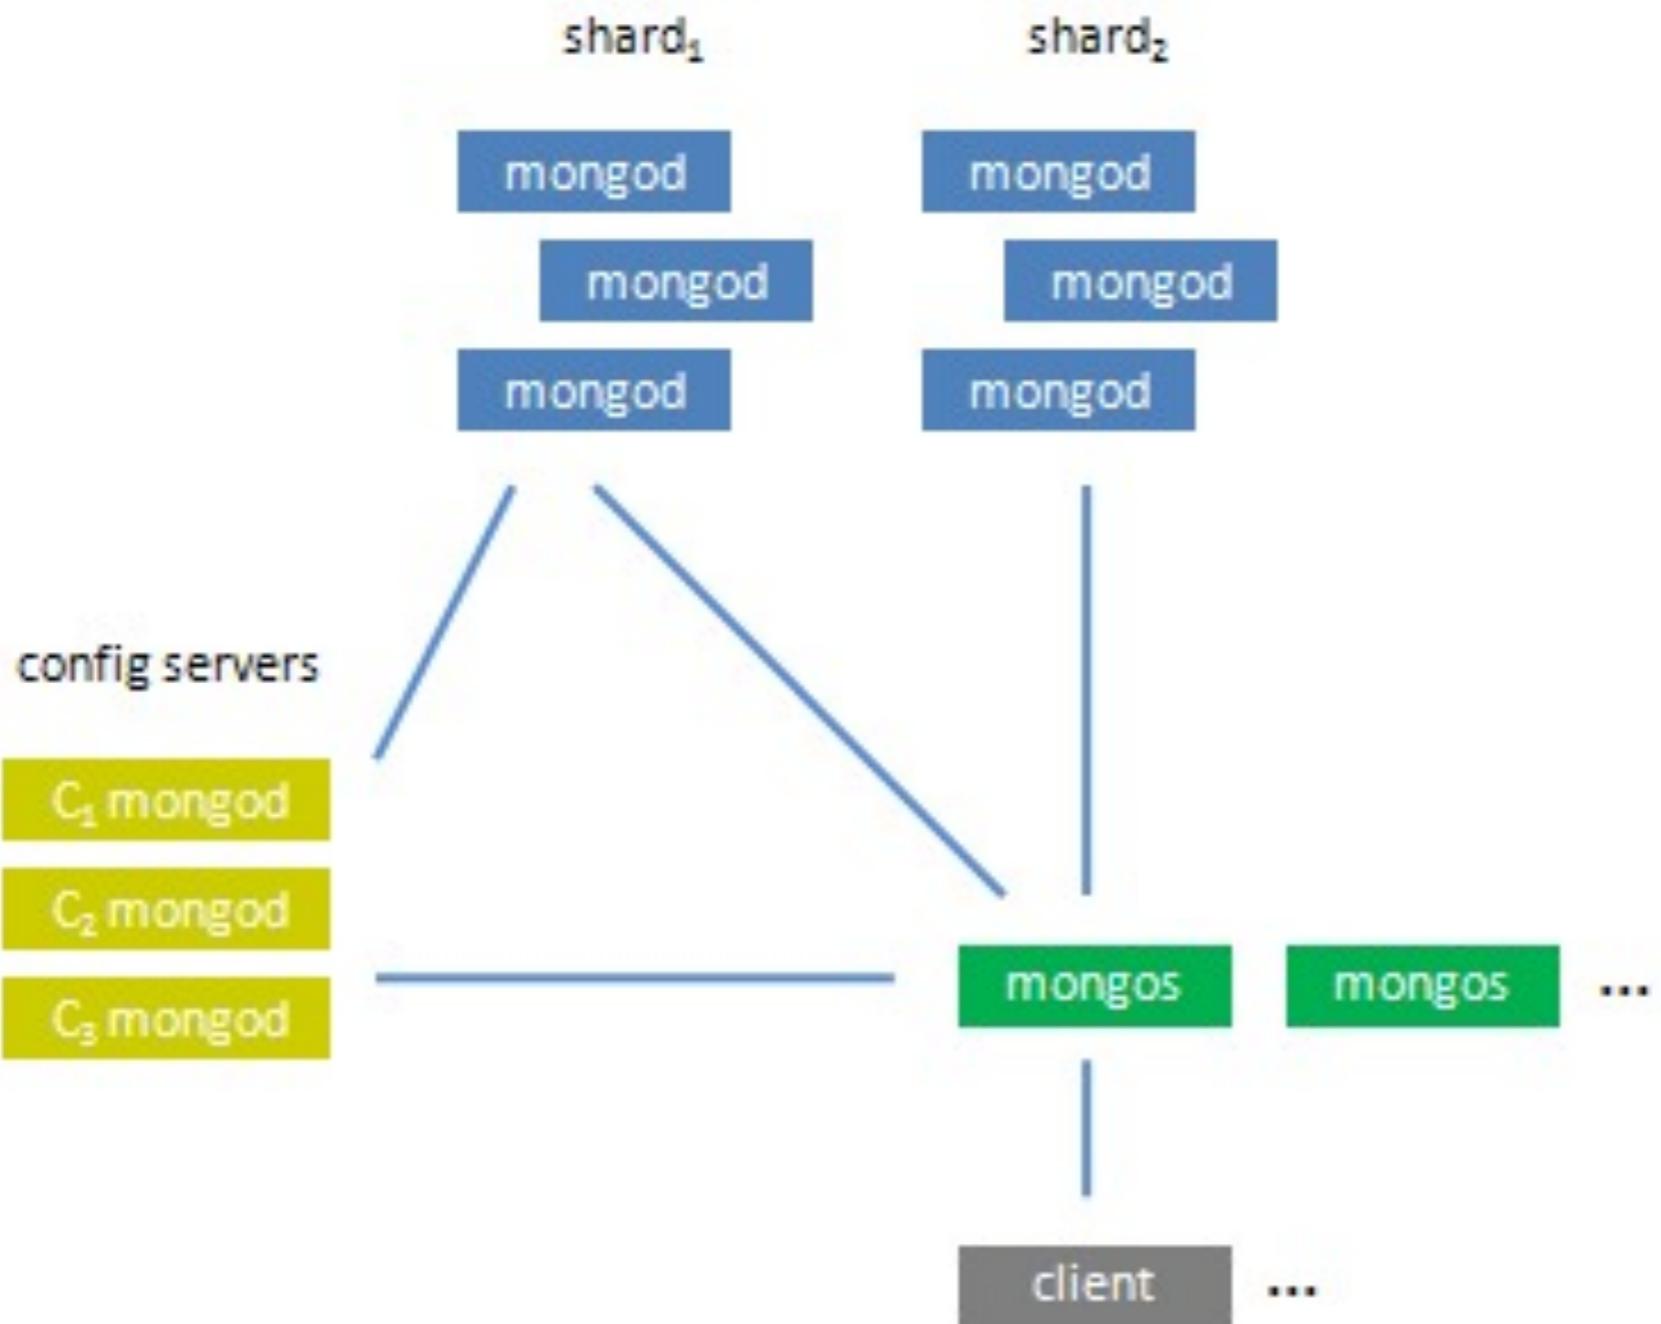
Typical scaling approaches don't work



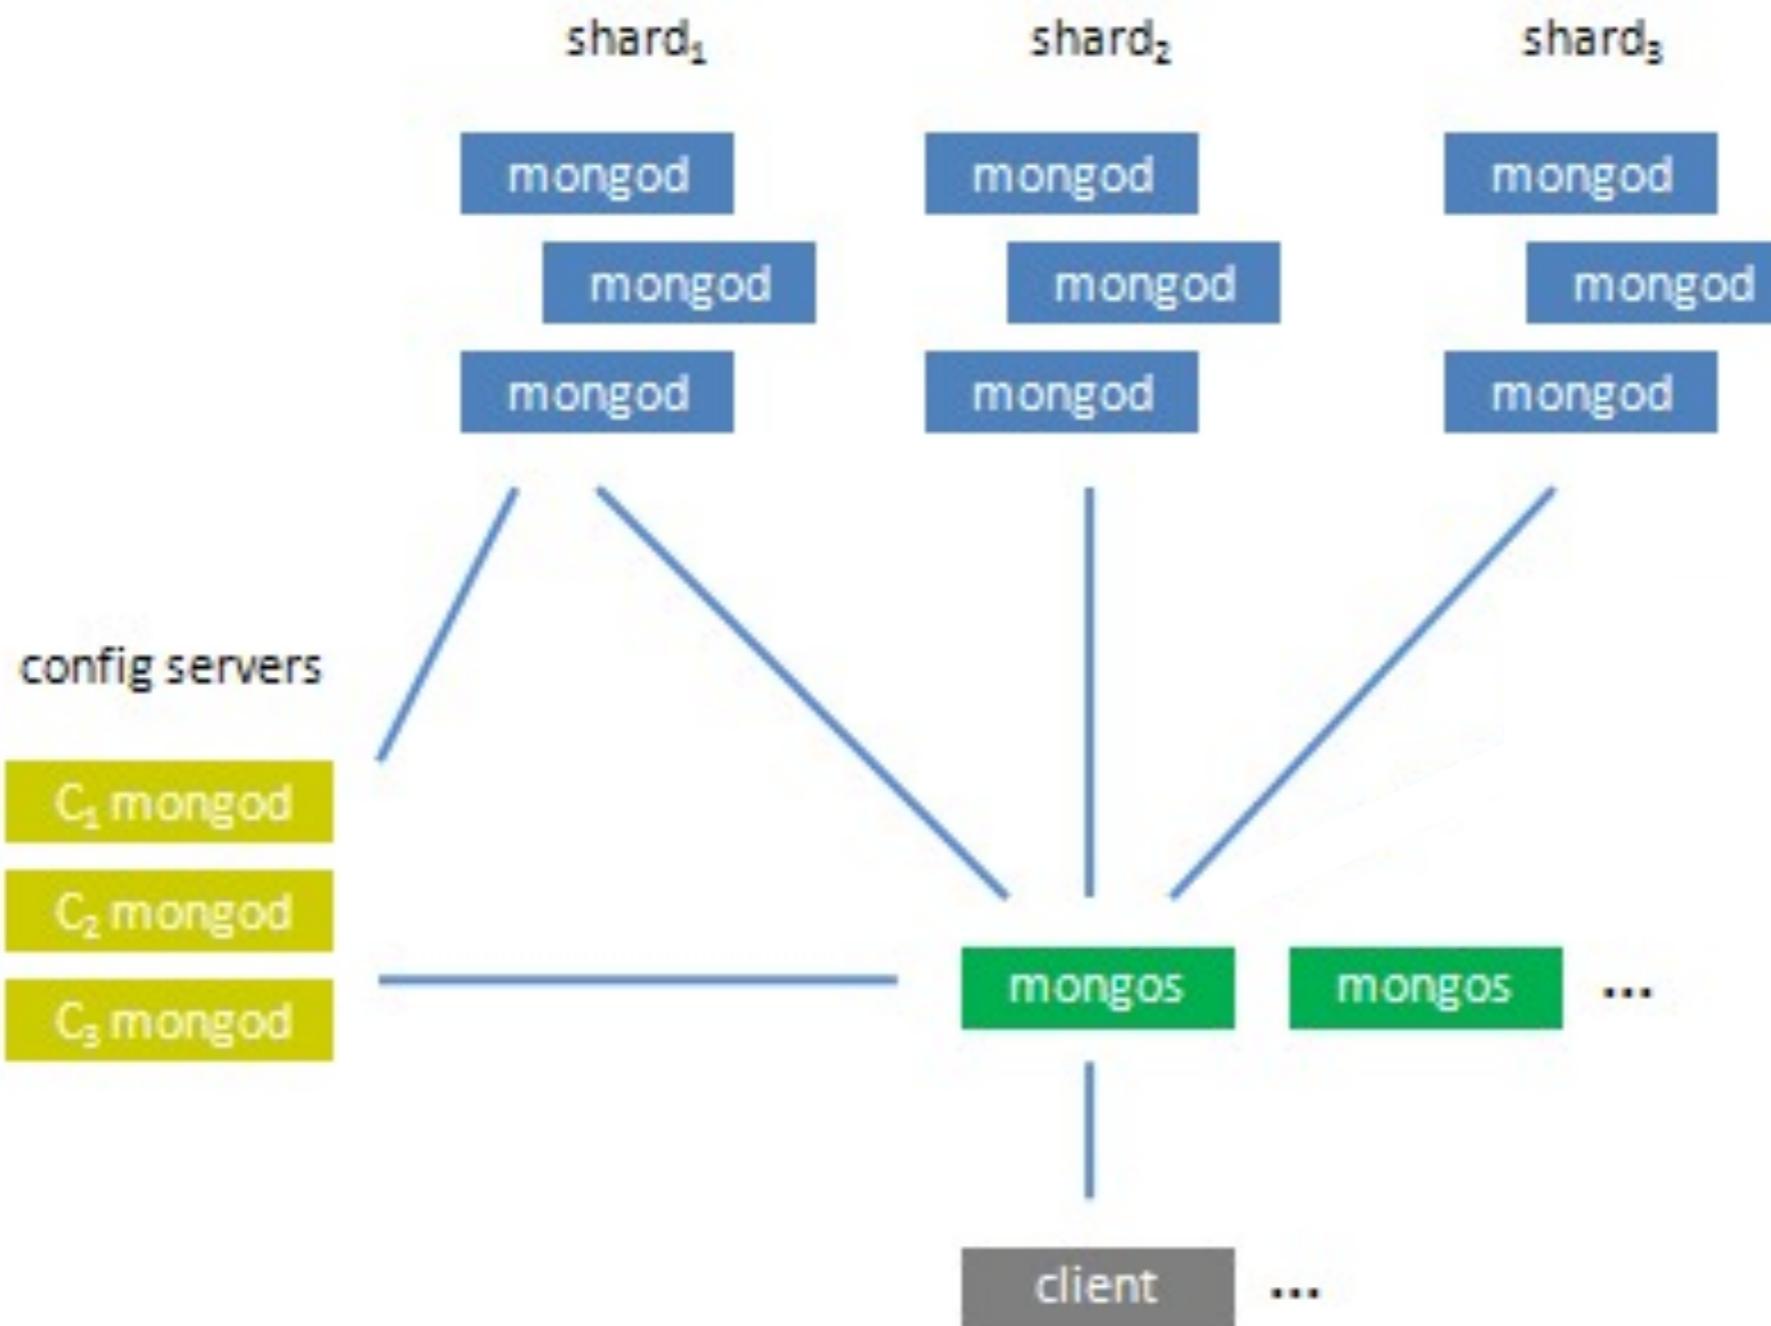
MongoDB scales out



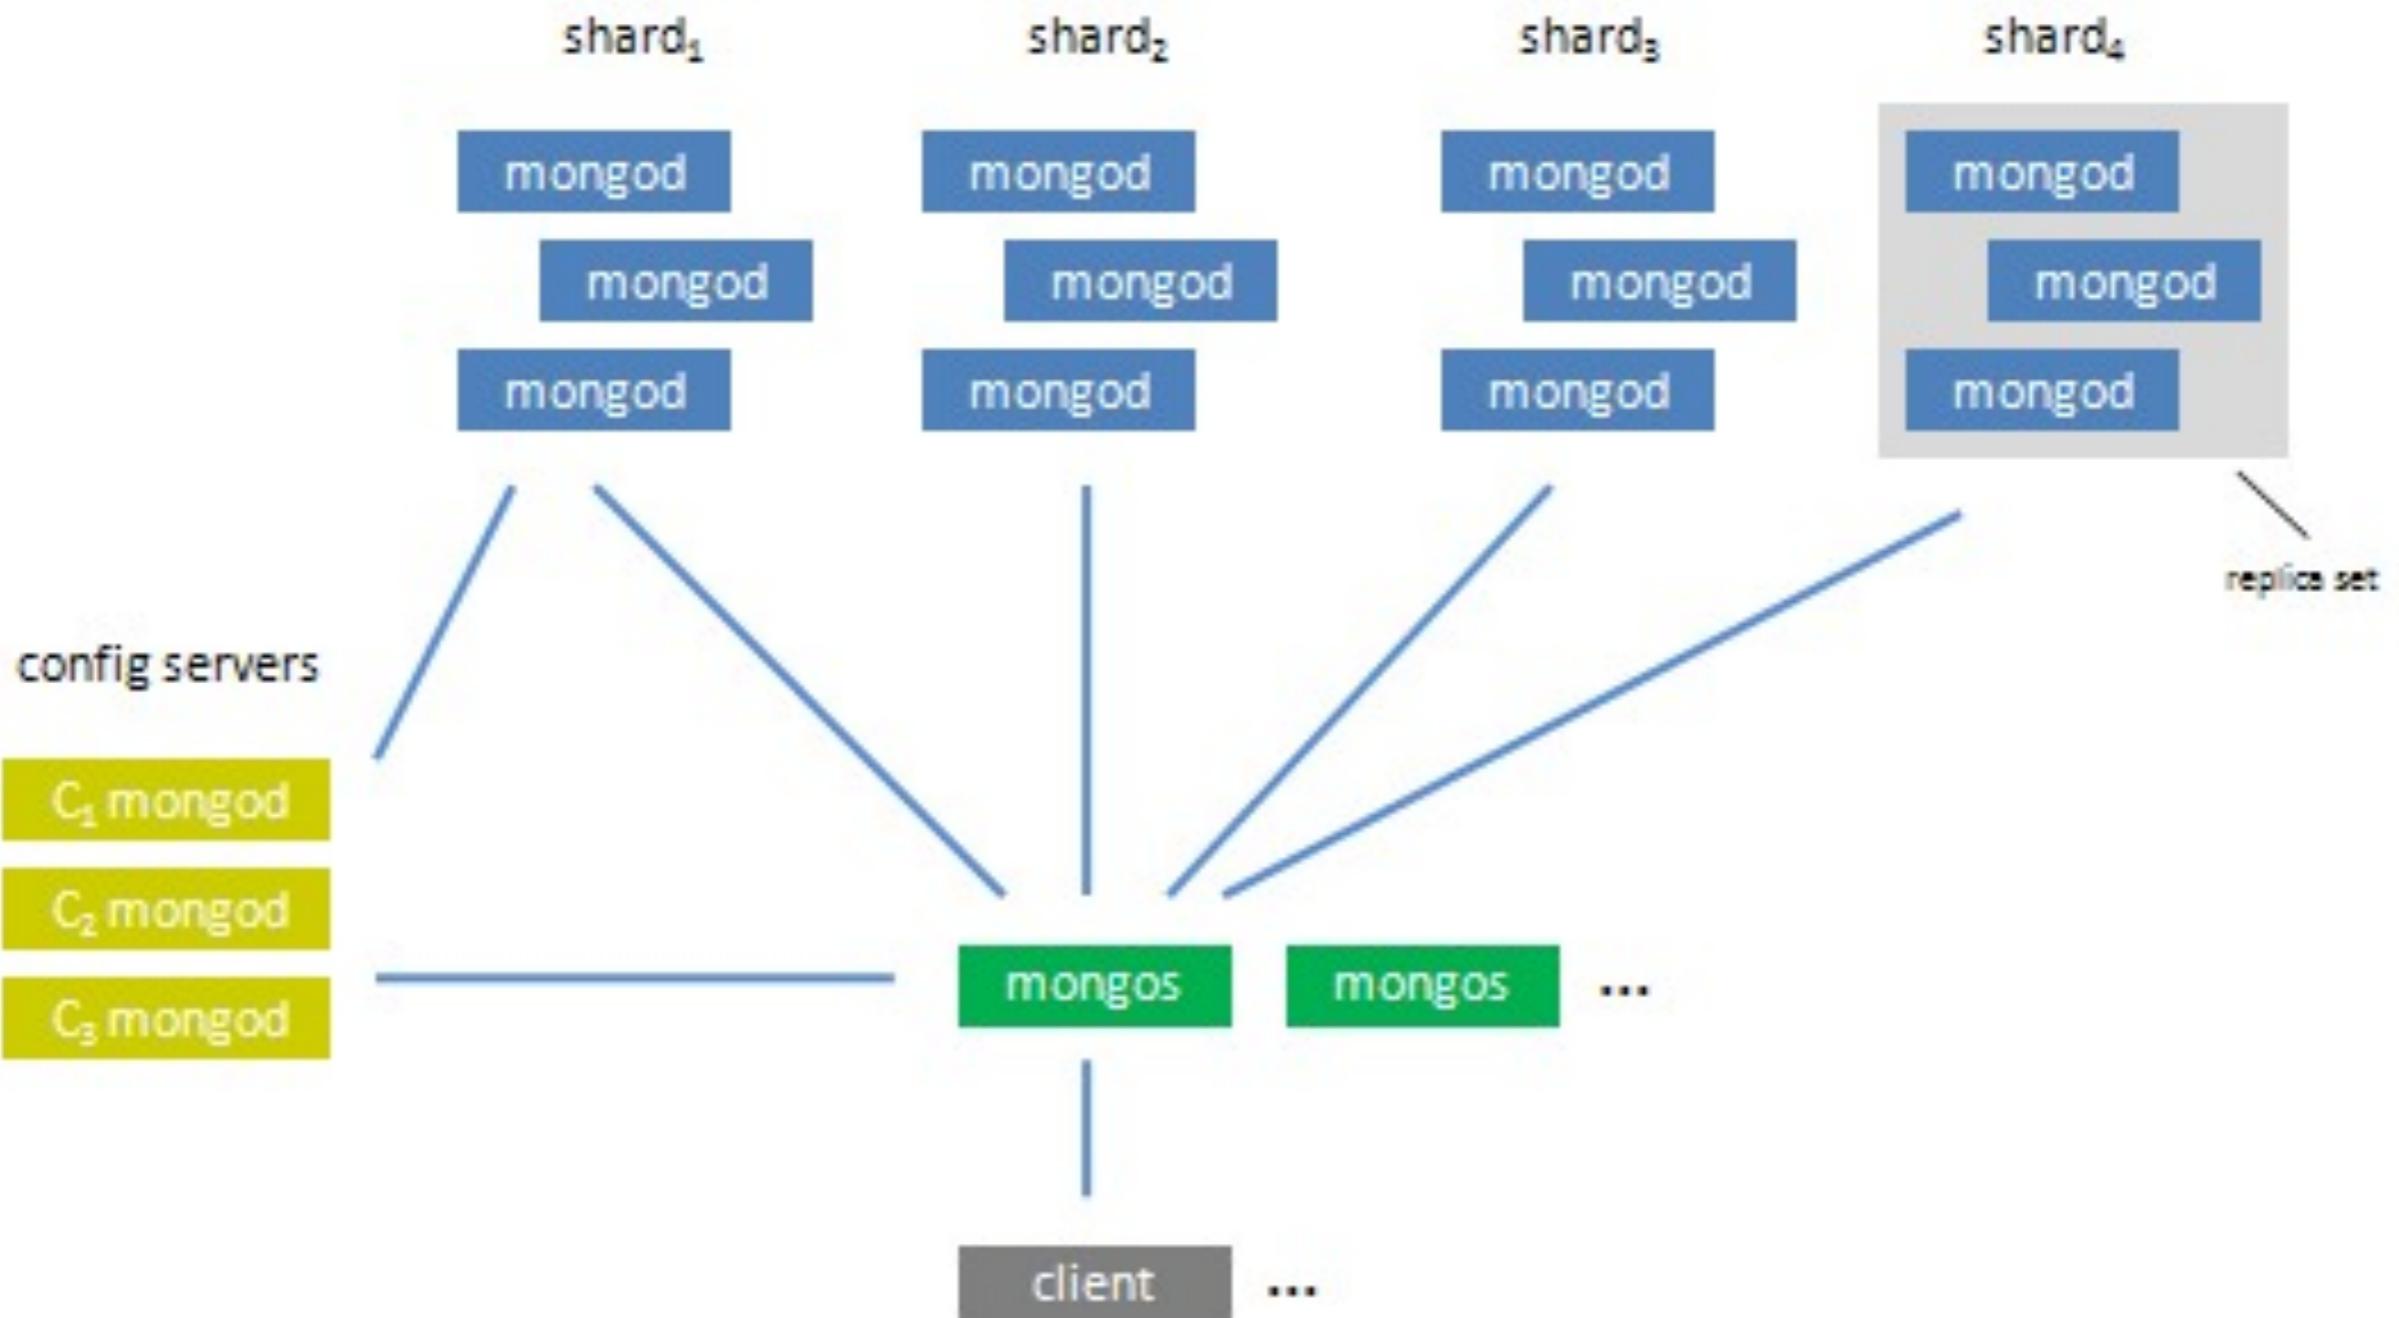
MongoDB scales out



MongoDB scales out



MongoDB scales out



MongoDB in the cloud

- Elastic database capacity
 - Add more storage, compute, memory dynamically
- Portable across environments
 - Public cloud
 - Private cloud
 - Micro cloud
- Simplifies operations
 - No need to engineer solutions for one-off hardware platforms
 - “Just another VM”

Q&A